

q2

November 25, 2018

```
In [1]: import numpy as np
import tensorflow as tf
from tensorflow.python.framework import ops
from tensorflow.python.ops import clip_ops
import matplotlib.pyplot as plt
from sklearn.datasets import load_digits
from sklearn.cross_validation import train_test_split

/anaconda3/envs/tfdeeplearning/lib/python3.5/site-packages/sklearn/cross_validation.py:41: DeprecationWarning:
  "This module will be removed in 0.20.", DeprecationWarning)

In [2]: #Loading MNIST dataset
X, y = load_digits(return_X_y=True)

In [3]: #Apply stratified sampling
X_train ,X_test,y_train,y_test = train_test_split(X,y,train_size = 0.7,random_state = 1)

In [4]: """Hyper-parameters"""
batch_size = 300                # Batch size for stochastic gradient descent
test_size = batch_size          # Temporary heuristic. In future we'd like to decouple test and train sizes
num_central = 150               # Number of "hidden neurons" that is number of centroids
max_iterations = 1000           # Max number of iterations
learning_rate = 5e-2            # Learning rate
num_classes = 10               # Number of target classes, 10 for MNIST
var_rbf = 225                  # What variance do you expect workable for the RBF?

#Obtain and proclaim sizes
N,D = X_train.shape
Ntest = X_test.shape[0]
print('We have %s observations with %s dimensions'%(N,D))

#Proclaim the epochs
epochs = np.floor(batch_size*max_iterations / N)
print('Train with approximately %d epochs' %(epochs))

We have 1257 observations with 64 dimensions
Train with approximately 238 epochs
```

```

In [5]: #Placeholders for data
        x = tf.placeholder('float',shape=[batch_size,D],name='input_data')
        y_ = tf.placeholder(tf.int64, shape=[batch_size], name = 'Ground_truth')

In [6]: with tf.name_scope("Hidden_layer") as scope:
        #Centroids and var are the main trainable parameters of the first layer
        centroids = tf.Variable(tf.random_uniform([num_cent, D], dtype=tf.float32), name='centroids')
        var = tf.Variable(tf.truncated_normal([num_cent, D], mean=0.0, stddev=0.1, dtype=tf.float32), name='var')

        #For now, we collect the distances
        exp_list = []
        for i in range(num_cent):
            exp_list.append(tf.exp((-1*tf.reduce_sum(tf.square(tf.subtract(x, centroids[i, :])), axis=1))))
        phi = tf.transpose(tf.stack(exp_list))

        with tf.name_scope("Output_layer") as scope:
            w = tf.Variable(tf.truncated_normal([num_cent, num_classes], stddev=0.1, dtype=tf.float32), name='weights')
            bias = tf.Variable(tf.constant(0.1, shape=[num_classes]), name='bias')

            h = tf.matmul(phi, w) + bias
            size2 = tf.shape(h)

            with tf.name_scope("Softmax") as scope:
                loss = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=h, labels=y_)
                cost = tf.reduce_sum(loss)
                loss_summ = tf.summary.scalar("cross_entropy_loss", cost)

INFO:tensorflow:Summary name cross_entropy_loss is illegal; using cross_entropy_loss instead.

In [7]: with tf.name_scope("train") as scope:
        tvars = tf.trainable_variables()
        #We clip the gradients to prevent explosion
        grads = tf.gradients(cost, tvars)
        optimizer = tf.train.AdamOptimizer(learning_rate)
        gradients = zip(grads, tvars)
        train_step = optimizer.apply_gradients(gradients)
        # The following block plots for every trainable variable
        num1 = tf.constant([[0]])
        for gradient, variable in gradients:
            if isinstance(gradient, ops.IndexedSlices):
                grad_values = gradient.values
            else:
                grad_values = gradient
            num1 += tf.reduce_sum(tf.size(variable))
            h1 = tf.histogram_summary(variable.name, variable)
            h2 = tf.histogram_summary(variable.name + "/gradients", grad_values)
            h3 = tf.histogram_summary(variable.name + "/gradient_norm", clip_ops.global_norm(grad_values))

```

```

with tf.name_scope("Evaluating") as scope:
    correct_prediction = tf.equal(tf.argmax(h,1), y_)
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
    accuracy_summary = tf.summary.scalar("accuracy", accuracy)

```

```
merged = tf.summary.merge_all()
```

```
perf_collect = np.zeros((4,int(np.floor(max_iterations /100))))
```

In [8]: *#Begin Training*

```

with tf.Session() as sess:
    with tf.device("/cpu:0"):
        print('Start session')
        writer = tf.summary.FileWriter("./log_tb", sess.graph_def)

        step = 0
        sess.run(tf.initialize_all_variables())

    for i in range(max_iterations):
        batch_ind = np.random.choice(N,batch_size,replace=False)
        if i%100 == 1:
            #Measure train performance
            result = sess.run([cost,accuracy,train_step],feed_dict={x:X_train[batch_ind], y: y_train[batch_ind]})
            perf_collect[0,step] = result[0]
            perf_collect[2,step] = result[1]

            #Measure test performance
            test_ind = np.random.choice(Ntest,test_size,replace=False)
            result = sess.run([cost,accuracy,merged], feed_dict={ x: X_test[test_ind], y: y_test[test_ind]})
            perf_collect[1,step] = result[0]
            perf_collect[3,step] = result[1]

            #Write information for Tensorboard
            summary_str = result[2]
            writer.add_summary(summary_str, i)
            writer.flush() #Don't forget this command! It makes sure Python writes the summaries

            #Print intermediate numbers to terminal
            acc = result[1]
            print("Estimated accuracy at iteration %s of %s: %s" % (i,max_iterations, acc))
            step += 1
        else:
            sess.run(train_step,feed_dict={x:X_train[batch_ind], y: y_train[batch_ind]})

```

Start session

WARNING:tensorflow:Passing a `GraphDef` to the SummaryWriter is deprecated. Pass a `Graph` object instead.


WARNING:tensorflow:From /anaconda3/envs/tfdeeplearning/lib/python3.5/site-packages/tensorflow/core/framework/summary_iterator.py:133: make_iterator (from tensorflow.python.framework.summary_iterator) is deprecated and will be removed in a future version. Use tf.nn.summary_iterator instead.

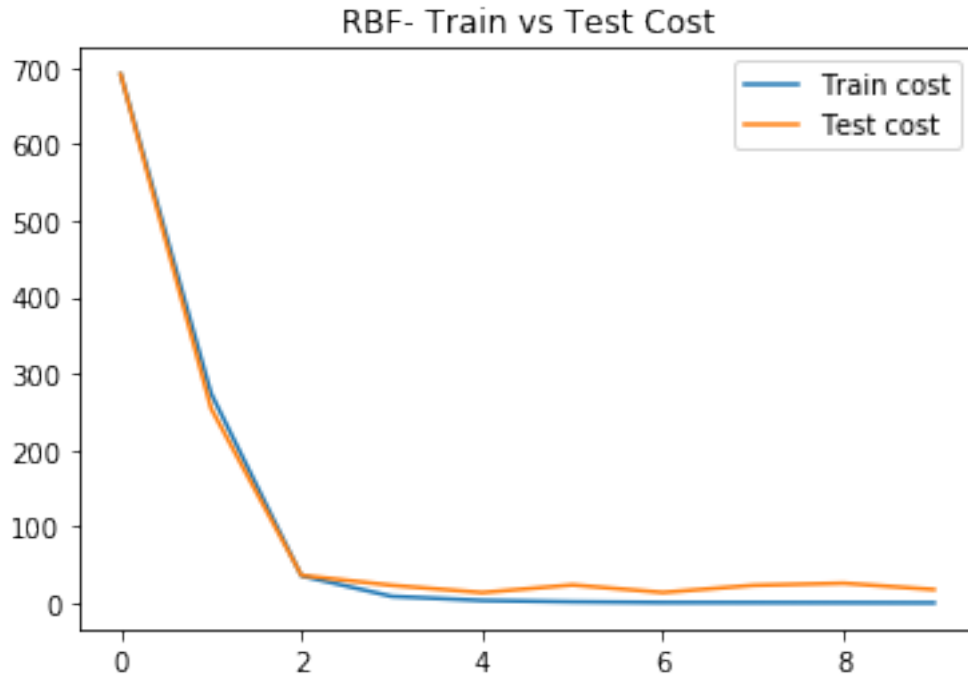
Instructions for updating:

Use ``tf.global_variables_initializer`` instead.

Estimated accuracy at iteration 1 of 1000: 0.0933333
Estimated accuracy at iteration 101 of 1000: 0.756667
Estimated accuracy at iteration 201 of 1000: 0.96
Estimated accuracy at iteration 301 of 1000: 0.973333
Estimated accuracy at iteration 401 of 1000: 0.983333
Estimated accuracy at iteration 501 of 1000: 0.966667
Estimated accuracy at iteration 601 of 1000: 0.983333
Estimated accuracy at iteration 701 of 1000: 0.97
Estimated accuracy at iteration 801 of 1000: 0.966667
Estimated accuracy at iteration 901 of 1000: 0.98

```
In [9]: """Additional plots"""  
plt.figure()  
plt.plot(perf_collect[2],label = 'Train accuracy')  
plt.plot(perf_collect[3],label = 'Test accuracy')  
plt.title('RBF- Train vs Test Accuracy')  
plt.legend()  
plt.show()  
  
plt.figure()  
plt.plot(perf_collect[0],label = 'Train cost')  
plt.plot(perf_collect[1],label = 'Test cost')  
plt.title('RBF- Train vs Test Cost')  
plt.legend()  
plt.show()
```





0.0.1 Changing the number of the hidden layer

```
In [10]: num_centra = 300
```

```
In [11]: with tf.name_scope("Hidden_layer") as scope:
    #Centroids and var are the main trainable parameters of the first layer
    centroids = tf.Variable(tf.random_uniform([num_centra,D],dtype=tf.float32),name='centroids')
    var = tf.Variable(tf.truncated_normal([num_centra],mean=var_rbf,stddev=5,dtype=tf.float32),name='var')

    #For now, we collect the distances
    exp_list = []
    for i in range(num_centra):
        exp_list.append(tf.exp((-1*tf.reduce_sum(tf.square(tf.subtract(x,centroids[i],var)),axis=1))))
    phi = tf.transpose(tf.stack(exp_list))

    with tf.name_scope("Output_layer") as scope:
        w = tf.Variable(tf.truncated_normal([num_centra,num_classes], stddev=0.1, dtype=tf.float32),name='weights')
        bias = tf.Variable(tf.constant(0.1, shape=[num_classes]),name='bias')

        h = tf.matmul(phi,w)+bias
        size2 = tf.shape(h)
```

```

with tf.name_scope("Softmax") as scope:
    loss = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=h,labels=y_)
    cost = tf.reduce_sum(loss)
    loss_summ = tf.summary.scalar("cross entropy_loss", cost)

```

INFO:tensorflow:Summary name cross entropy_loss is illegal; using cross_entropy_loss instead.

```

In [12]: with tf.name_scope("train") as scope:
    tvars = tf.trainable_variables()
    #We clip the gradients to prevent explosion
    grads = tf.gradients(cost, tvars)
    optimizer = tf.train.AdamOptimizer(learning_rate)
    gradients = zip(grads, tvars)
    train_step = optimizer.apply_gradients(gradients)
    # The following block plots for every trainable variable

    numel = tf.constant([[0]])
    for gradient, variable in gradients:
        if isinstance(gradient, ops.IndexedSlices):
            grad_values = gradient.values
        else:
            grad_values = gradient
        numel += tf.reduce_sum(tf.size(variable))
        h1 = tf.histogram_summary(variable.name, variable)
        h2 = tf.histogram_summary(variable.name + "/gradients", grad_values)
        h3 = tf.histogram_summary(variable.name + "/gradient_norm", clip_ops.global_norm(grad_values))
    with tf.name_scope("Evaluating") as scope:
        correct_prediction = tf.equal(tf.argmax(h,1), y_)
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
        accuracy_summary = tf.summary.scalar("accuracy", accuracy)

merged = tf.summary.merge_all()

# For now, we collect performances in a Numpy array.
# In future releases, I hope TensorBoard allows for more
# flexibility in plotting
perf_collect = np.zeros((4,int(np.floor(max_iterations /100))))

```

```

In [13]: #Begin Training
with tf.Session() as sess:
    with tf.device("/cpu:0"):
        print('Start session')
        writer = tf.summary.FileWriter("./log_tb", sess.graph_def)

        step = 0
        sess.run(tf.initialize_all_variables())

```

```

# #Debugging
# batch_ind = np.random.choice(N,batch_size,replace=False)
# result = sess.run([phi],feed_dict={x:X_train[batch_ind], y_: y_train[batch_ind]})
# print(result[0])

for i in range(max_iterations):
    batch_ind = np.random.choice(N,batch_size,replace=False)
    if i%100 == 1:
        #Measure train performance
        result = sess.run([cost,accuracy,train_step],feed_dict={x:X_train[batch_ind], y_: y_train[batch_ind]})
        perf_collect[0,step] = result[0]
        perf_collect[2,step] = result[1]

        #Measure test performance
        test_ind = np.random.choice(Ntest,test_size,replace=False)
        result = sess.run([cost,accuracy,merged], feed_dict={ x: X_test[test_ind], y_: y_test[test_ind]})
        perf_collect[1,step] = result[0]
        perf_collect[3,step] = result[1]

        #Write information for Tensorboard
        summary_str = result[2]
        writer.add_summary(summary_str, i)
        writer.flush() #Don't forget this command! It makes sure Python writes to disk

        #Print intermediate numbers to terminal
        acc = result[1]
        print("Estimated accuracy at iteration %s of %s: %s" % (i,max_iterations, acc))
        step += 1
    else:
        sess.run(train_step,feed_dict={x:X_train[batch_ind], y_: y_train[batch_ind]})

```

Start session

WARNING:tensorflow:Passing a `GraphDef` to the SummaryWriter is deprecated. Pass a `Graph` object instead.

WARNING:tensorflow:From /anaconda3/envs/tfdeeplearning/lib/python3.5/site-packages/tensorflow/core/framework/summary_iterator.py:133: make_iterator (from tensorflow.python.framework.summary_iterator) is deprecated and will be removed in a future version. Use tf.nn.summary_iterator instead.

Instructions for updating:

Use `tf.nn.summary_iterator` instead.

Estimated accuracy at iteration 1 of 1000: 0.106667

Estimated accuracy at iteration 101 of 1000: 0.85

Estimated accuracy at iteration 201 of 1000: 0.95

Estimated accuracy at iteration 301 of 1000: 0.973333

Estimated accuracy at iteration 401 of 1000: 0.97

Estimated accuracy at iteration 501 of 1000: 0.966667

Estimated accuracy at iteration 601 of 1000: 0.976667

Estimated accuracy at iteration 701 of 1000: 0.973333

Estimated accuracy at iteration 801 of 1000: 0.986667

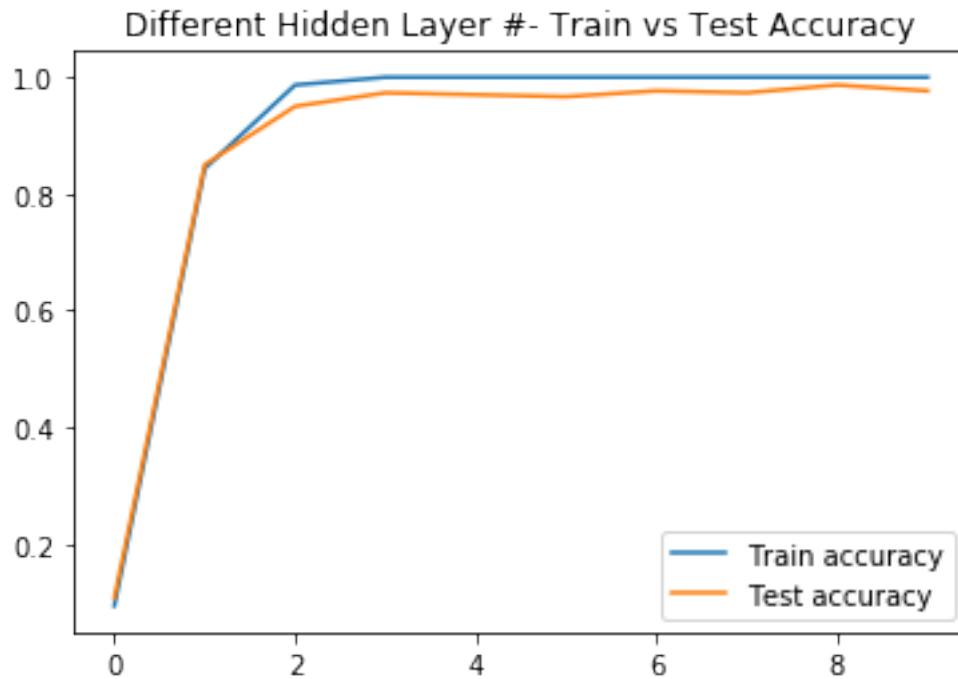
Estimated accuracy at iteration 901 of 1000: 0.976667


```

In [14]: """Additional plots"""
plt.figure()
plt.plot(perf_collect[2],label = 'Train accuracy')
plt.plot(perf_collect[3],label = 'Test accuracy')
plt.title('Different Hidden Layer #- Train vs Test Accuracy')
plt.legend()
plt.show()

plt.figure()
plt.plot(perf_collect[0],label = 'Train cost')
plt.plot(perf_collect[1],label = 'Test cost')
plt.title('Different Hidden Layer #- Train vs Test Cost')
plt.legend()
plt.show()

```





0.0.2 Adding a dropout

```
In [15]: with tf.name_scope("Hidden_layer") as scope:
    #Centroids and var are the main trainable parameters of the first layer
    centroids = tf.Variable(tf.random_uniform([num_cent, D], dtype=tf.float32), name='c')
    var = tf.Variable(tf.truncated_normal([num_cent], mean=var_rbf, stddev=5, dtype=tf.float32), name='var')

    #For now, we collect the distances
    exp_list = []
    for i in range(num_cent):
        exp_list.append(tf.exp((-1*tf.reduce_sum(tf.square(tf.subtract(x, centroids[i, :])), axis=1))
        phi = tf.transpose(tf.stack(exp_list))

    with tf.name_scope("Output_layer") as scope:
        w = tf.Variable(tf.truncated_normal([num_cent, num_classes], stddev=0.1, dtype=tf.float32), name='w')
        bias = tf.Variable(tf.constant(0.1, shape=[num_classes]), name='bias')

        h = tf.matmul(phi, w) + bias
        #Adding dropout
        h = tf.nn.dropout(h, 0.5)

        size2 = tf.shape(h)

    with tf.name_scope("Softmax") as scope:
        loss = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=h, labels=y_)
```

```
cost = tf.reduce_sum(loss)
loss_summ = tf.summary.scalar("cross_entropy_loss", cost)
```

INFO:tensorflow:Summary name cross_entropy_loss is illegal; using cross_entropy_loss instead.

```
In [16]: with tf.name_scope("train") as scope:
    tvars = tf.trainable_variables()
    #We clip the gradients to prevent explosion
    grads = tf.gradients(cost, tvars)
    optimizer = tf.train.AdamOptimizer(learning_rate)
    gradients = zip(grads, tvars)
    train_step = optimizer.apply_gradients(gradients)
    # The following block plots for every trainable variable

    numel = tf.constant([[0]])
    for gradient, variable in gradients:
        if isinstance(gradient, ops.IndexedSlices):
            grad_values = gradient.values
        else:
            grad_values = gradient
        numel += tf.reduce_sum(tf.size(variable))
        h1 = tf.histogram_summary(variable.name, variable)
        h2 = tf.histogram_summary(variable.name + "/gradients", grad_values)
        h3 = tf.histogram_summary(variable.name + "/gradient_norm", clip_ops.global_norm(grad_values))
    with tf.name_scope("Evaluating") as scope:
        correct_prediction = tf.equal(tf.argmax(h,1), y_)
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
        accuracy_summary = tf.summary.scalar("accuracy", accuracy)

    merged = tf.summary.merge_all()

    perf_collect = np.zeros((4,int(np.floor(max_iterations /100))))
```

```
In [17]: #Begin Training
    with tf.Session() as sess:
        with tf.device("/cpu:0"):
            print('Start session')
            writer = tf.summary.FileWriter("./log_tb", sess.graph_def)

            step = 0
            sess.run(tf.initialize_all_variables())

            for i in range(max_iterations):
                batch_ind = np.random.choice(N, batch_size, replace=False)
                if i%100 == 1:
                    #Measure train performance
                    result = sess.run([cost, accuracy, train_step], feed_dict={x:X_train[batch_ind], y:Y_train[batch_ind]})
```

```

perf_collect[0,step] = result[0]
perf_collect[2,step] = result[1]

#Measure test performance
test_ind = np.random.choice(Ntest,test_size,replace=False)
result = sess.run([cost,accuracy,merged], feed_dict={ x: X_test[test_ind]
perf_collect[1,step] = result[0]
perf_collect[3,step] = result[1]

#Write information for Tensorboard
summary_str = result[2]
writer.add_summary(summary_str, i)
writer.flush() #Don't forget this command! It makes sure Python writes t

#Print intermediate numbers to terminal
acc = result[1]
print("Estimated accuracy at iteration %s of %s: %s" % (i,max_iterations,
step += 1
else:
    sess.run(train_step,feed_dict={x:X_train[batch_ind], y_: y_train[batch_in

```

Start session

WARNING:tensorflow:Passing a `GraphDef` to the SummaryWriter is deprecated. Pass a `Graph` obj

WARNING:tensorflow:From /anaconda3/envs/tfdeeplearning/lib/python3.5/site-packages/tensorflow/

Instructions for updating:

Use `tf.global_variables_initializer` instead.

Estimated accuracy at iteration 1 of 1000: 0.106667

Estimated accuracy at iteration 101 of 1000: 0.5

Estimated accuracy at iteration 201 of 1000: 0.503333

Estimated accuracy at iteration 301 of 1000: 0.563333

Estimated accuracy at iteration 401 of 1000: 0.546667

Estimated accuracy at iteration 501 of 1000: 0.59

Estimated accuracy at iteration 601 of 1000: 0.613333

Estimated accuracy at iteration 701 of 1000: 0.606667

Estimated accuracy at iteration 801 of 1000: 0.55

Estimated accuracy at iteration 901 of 1000: 0.633333

In [18]: *"""Additional plots"""*

```

plt.figure()
plt.plot(perf_collect[2],label = 'Train accuracy')
plt.plot(perf_collect[3],label = 'Test accuracy')
plt.legend()
plt.title('Dropout- Train vs Test Accuracy')
plt.show()

```

```

plt.figure()
plt.plot(perf_collect[0],label = 'Train cost')

```

```
plt.plot(perf_collect[1],label = 'Test cost')
plt.legend()
plt.title('Dropout- Train vs Test Cost')
plt.show()
```

